

# **Interview Assignment:**

## **Intelligent message routing with guardrails**

Code Repo: [github.com/Franklin112233/intelligent\\_message\\_routing](https://github.com/Franklin112233/intelligent_message_routing)

Linlin Yu

Feb 2026

# 1. Goal, features, and KPI

---

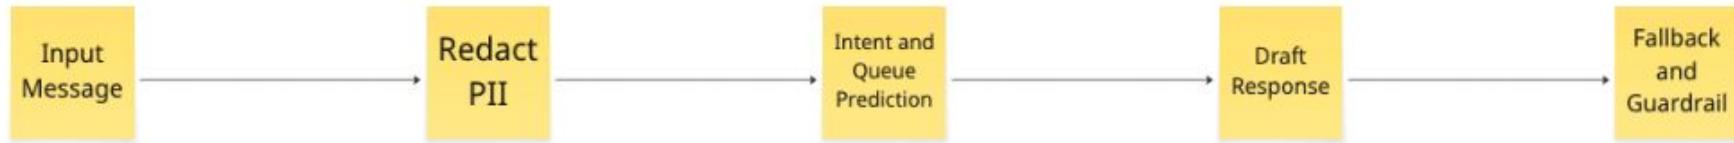
## Goal

Triage customer messages: redact PII → classify intent → policy-grounded draft with guardrails and fallbacks.

## Main features

| Feature               | What it does  | KPI  |
|-----------------------|---|--|
| PII redaction         | Redact sensitive tokens before any external call (YAML patterns).                 | Redaction recall; privacy/compliance.                  |
| Intent classification | Route to queue (Fraud, Disputes, General, Credit); Multi-Task Learning (MTL).     | Precision (e.g. fraud routing); latency SLA.           |
| Draft response        | Policy-grounded draft for $\geq 2$ intents; template or LLM; confidence fallback. | Draft quality; citation; cost (fallback saves tokens). |
| Guardrails            | Citation + PII-in-draft checks; escalate on low confidence or fail.               | Safe automation; audit trail.                          |
| Evaluation            | Classification metrics, holdout eval, redaction/draft tests.                      | Accuracy; redaction recall.                            |

## 2. App runnable example



Enter message (or press Enter to run 5 from CSV) (): Someone set up a new payee I don't know. My postcode is W1A 0AX

Config

Backend: mtl Draft: LLM (GPT-4o-mini)

Input Message

Someone set up a new payee I don't know. My postcode is W1A 0AX

Redaction

Someone set up a new payee I don't know. My postcode is [POSTCODE]

Intent/Queue Prediction

Intent: fraud  
Queue: Fraud/Economic Crime Prevention  
Confidence: 0.83  
Fallback: False  
Checks: OK

Draft

Thank you for reaching out. I recommend that we immediately freeze your card to prevent any unauthorized transactions. Please change your credentials as well, and I will begin a fraud investigation to confirm recent activity. You can expect updates within 7-10 business days regarding the next steps .

### 3.Solution design

#### . Hybrid model: LLM + MTL(Multi-Task Learning)

| Component | Solves  | Role  |
|-----------|---|---|
| MTL       | Intent + queue; latency & cost; explainability; data local.               | TF-IDF + 2 LogReg; Inference inside the interal environment; no raw data out. |
| LLM       | Workflow orchestrator. Draft text; policy + citations; flexible phrasing. | Draft only (redacted + intent); fallback; guardrails checks.                  |

Benefits: fast MTL train; task tuning; local/cheap/secure; fallback cuts cost; routing + cited replies.

#### . Workflow integration

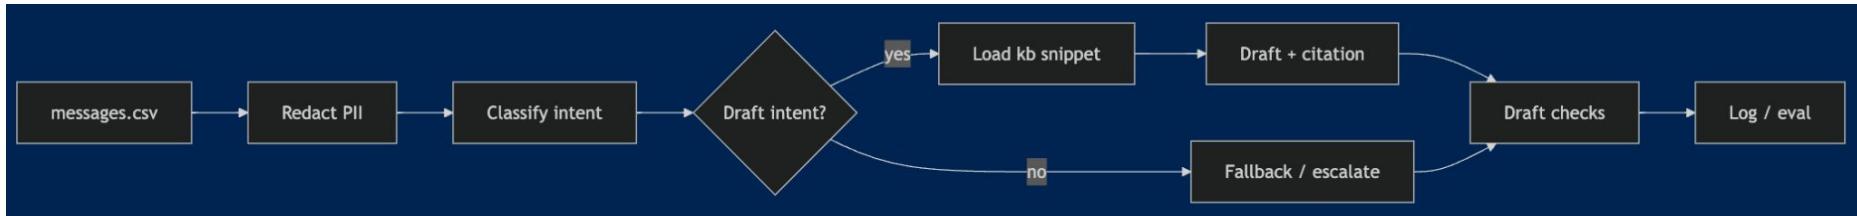
Redact → Model (MTL) → Draft (kb → template or LLM) → Guardrails (citation, PII-in-draft; escalate). Redaction gates all; MTL picks policy; guardrails on final draft.

#### . Spec driven development

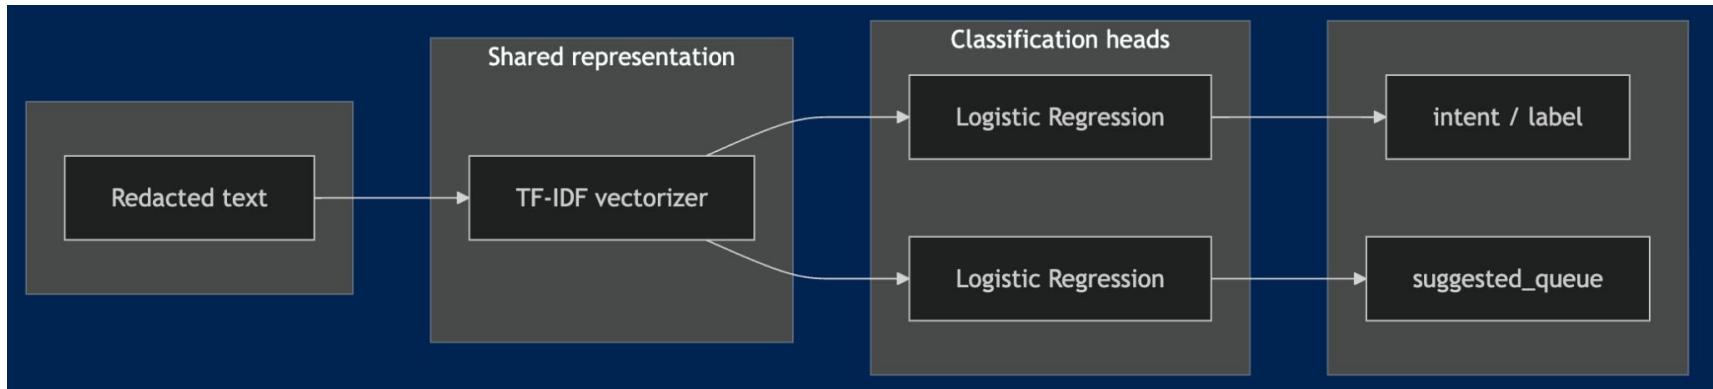
Proposal → design → specs → tasks -> changes; single source of truth; auditable; versioned.

## 4. Architecture

.System architecture (pipeline)



.MTL structure (shared representation, two heads)



## 5. Evaluation, explainability & monitoring

| Area                          | What we do / key points  |
|-------------------------------|--|
| Evaluation                    | Classification accuracy (optional holdout); redaction unit tests + recall; draft guardrail pass rate and fallback rate on sample.  |
| Explainability                | MTL = features/weights, auditable; LLM = citations + guardrails; escalation on low confidence.   |
| Monitoring                    | Guardrail fail rate, fallback rate, latency; accuracy on labelled sample; <b>drift detection</b> (prediction vs ground truth over time → alert, retrain or investigate). |
| Testing safely pre-production | Shadow mode (no customer send); holdout + guardrails pass before go-live; human review sample; rollout with template/high-confidence first                               |

## 6. Implemented vs future improvements

| Area                  | Implemented  | Future  |
|-----------------------|--|---|
| PII redaction         | YAML + regex, tests; before external call.         | Rules-based metrics; more patterns.                               |
| Intent classification | MTL (TF-IDF + 2 LogReg); stub; holdout.            | More intents; LLM few-shot; Bert Tuning for semantic information. |
| Draft                 | ≥2 intents; template/LLM; 0.7, fallback.           | More intents; Internal LLM.                                       |
| Guardrails            | Citation + PII-in-draft; escalate.                 | Safety filter; grammar; thresholds; Business rules/compliance.    |
| Evaluation            | Metrics; draft checks; redaction tests; holdout.   | Recall curve; per-intent; Fairness.                               |
| CLI / run             | Interactive; panels                                | API; batch.   |
| Ops / monitoring      | Logs, eval.  | Dashboards; drift; alerting; retrain.                             |
| Risk                  | Redact first; no PII out; escalation.              | Risk register; audit; data/model governance                       |
| Cost                  | MTL cheap; LLM per-request; fallback saves tokens. | Alerts; quota; cost per queue; on-prem LLM                        |
| Latency               | MTL ms; LLM 100s ms; fixed order.                  | SLA; async; quantization.   |